

Projet de compilation avancée

Vérification statique des séquences d'appels aux fonctions collectives MPI

Hugo Lallemant, Mickaël Saes-Vincensini

5 Novembre 2024

Contexte

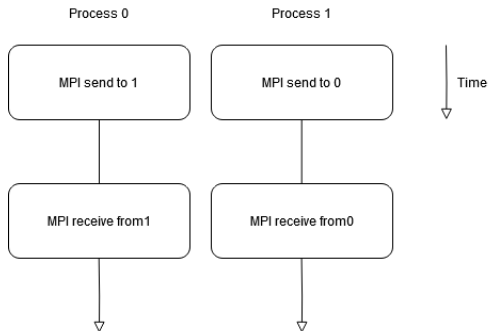





Schéma d'un deadlock dans une exécution de code

Préparation du CFG



```
typedef struct mpi_ranks
{
    int code;
    int * ranks;
} mpi_ranks;
```

Modification du CFG




Rang des collectives

$$\text{Rang}(n) = \max_{x \in \text{preds}(n)} \text{Rang}(x)$$

Équation que nous résolvons pour calculer le rang des collectives

```
index: 0 - collective: 5 - [0, 0, 0, 0, 0, ]
index: 2 - collective: 0 - [1, 0, 0, 0, 0, ]
index: 3 - collective: 4 - [1, 0, 0, 0, 1, ]
index: 4 - collective: 4 - [1, 0, 0, 0, 2, ]
index: 13 - collective: 4 - [1, 0, 0, 0, 3, ]
index: 5 - collective: 5 - [1, 0, 0, 0, 3, ]
index: 6 - collective: 5 - [1, 0, 0, 0, 3, ]
index: 7 - collective: 5 - [1, 0, 0, 0, 3, ]
index: 8 - collective: 4 - [1, 0, 0, 0, 2, ]
index: 9 - collective: 5 - [1, 0, 0, 0, 3, ]
index: 10 - collective: 5 - [1, 0, 0, 0, 0, ]
index: 11 - collective: 1 - [1, 1, 0, 0, 0, ]
index: 12 - collective: 5 - [1, 1, 0, 0, 3, ]
index: 1 - collective: 5 - [1, 1, 0, 0, 3, ]
code: 0, rank: 1 - 2
code: 1, rank: 1 - 11
code: 4, rank: 1 - 3
code: 4, rank: 2 - 4, 8
code: 4, rank: 3 - 13
```




Post dominance des ensembles

```

---- set postdominated ----
code: 0, rank: 1 - 2
code: 1, rank: 1 - 11
code: 4, rank: 1 - 3
code: 4, rank: 2 - 3, 4, 8
code: 4, rank: 3 - 4, 13


code: 0, rank: 1 - 2
code: 1, rank: 1 - 11
code: 4, rank: 1 - 3
code: 4, rank: 2 - 4, 8
code: 4, rank: 3 - 13

```



Frontière de post dominance des noeud

```
----- post dominance frontier -----
Basic Block: 0 :
Basic Block: 2 :
Basic Block: 3 : 10
Basic Block: 4 : 3
Basic Block: 13 : 3
Basic Block: 5 : 13
Basic Block: 6 : 13
Basic Block: 7 : 6
Basic Block: 8 : 3
Basic Block: 9 : 6, 13
Basic Block: 10 : 6, 13
Basic Block: 11 : 10
Basic Block: 12 :
Basic Block: 1 :
-----
```



Frontière de post dominance des ensembles

Soit $PDF(E) = \{ X \mid \exists S \in Succs(X) \text{ tq } E \gg_{sp} S, E \not\gg_p X \}$

Soit $F = \{ X \mid \exists n \in E \text{ tq } E \gg_{sp} n, X \in PDF(n), E \not\gg_p X \}$

On montre que $E=F$

```

---- set frontiers ----
code: 0, rank: 1 -
code: 1, rank: 1 - 10
Potential MPI Deadlock
code: 4, rank: 1 - 10
Potential MPI Deadlock
code: 4, rank: 2 - 10
Potential MPI Deadlock
code: 4, rank: 3 - 3
Potential MPI Deadlock

code: 0, rank: 1 - 2
code: 1, rank: 1 - 11
code: 4, rank: 1 - 3
code: 4, rank: 2 - 4, 8
code: 4, rank: 3 - 13
    
```




Frontière itérée

---- set iterated frontiers ----

```
code: 0, rank: 1 -
code: 1, rank: 1 - 3, 6, 10, 13
Potential MPI Deadlock
code: 4, rank: 1 - 3, 6, 10, 13
Potential MPI Deadlock
code: 4, rank: 2 - 3, 6, 10, 13
Potential MPI Deadlock
code: 4, rank: 3 - 3, 6, 10, 13
Potential MPI Deadlock
```

```
code: 0, rank: 1 - 2
code: 1, rank: 1 - 11
code: 4, rank: 1 - 3
code: 4, rank: 2 - 4, 8
code: 4, rank: 3 - 13
```



Affichage d'un warning

```
[mickael.saes-vincensini@hpc01 coav-2024]$ make test6
mkdir -p bin
g++ 1220 -I'gcc_1220 -print-file-name=plugin'/include -g -Wall -fno-rtti -shared -fPIC -o bin/libplugin.so src/mpi_plugin.cpp
mpicc tests/test6.c -g -O3 -o bin/test6 -fplugin=./bin/libplugin.so
plugin_init: Entering...
plugin_init: Check ok...
plugin_init: Pass added...
Now starting to examine function mpi_invalid
tests/test6.c: In function 'mpi_invalid':
tests/test6.c:18:9: warning: Potential issue: MPI collective MPI_Barrier in block 6
   18 |         MPI_Barrier(MPI_COMM_WORLD);
      |         ^
tests/test6.c:13:12: warning: Potential issue caused by the following fork in block 2
   13 |         if (c > 5) {
      |         ^
Now starting to examine function mpi_valid
No potential deadlock found.
Now starting to examine function main
tests/test6.c: In function 'main':
tests/test6.c:38:7: warning: Potential issue: MPI collective MPI_Barrier in block 4
   38 |         MPI_Barrier(MPI_COMM_WORLD);
      |         ^
tests/test6.c:35:6: warning: Potential issue caused by the following fork in block 2
   35 |         if (c < 20) {
      |         ^
tests/test6.c:36:8: warning: Potential issue caused by the following fork in block 3
   36 |         if (c < 10)
      |         ^
tests/test6.c:40:5: warning: Potential issue: MPI collective MPI_Barrier in block 5
   40 |         MPI_Barrier(MPI_COMM_WORLD);
      |         ^
tests/test6.c:35:6: warning: Potential issue caused by the following fork in block 2
   35 |         if (c < 20) {
      |         ^
```

Figure – Affichage généré par le warning

Gestion des pragma

```
#pragma ProjetCA mpicoll_check f  
#pragma ProjetCA mpicoll_check (f1, f2, ..., fn)
```

Conclusion et perspectives

```
unsigned int execute (function *fun)
{
    cfgviz_dump(fun, "initial");
    prepare_cfg(fun);
    cfgviz_dump(fun, "split");
    calculate_dominance_info(CDI_POST_DOMINATORS);
    bitmap_head *frontiers = post_dominance_frontiers(fun);
    bitmap_head *invalid_edges = cfg_prune(fun);
    cfgviz_dump(fun, "invalid_edges", invalid_edges);
    calculate_rank(fun, invalid_edges);
    bitmap_head **sets = collective_rank_set(fun);
    bitmap_head **set_postdominated = set_post_dominance(fun, sets);
    bitmap_head **set_frontiers = set_post_dominance_frontiers(fun, sets,
                                                                set_postdominated, frontiers);
    bitmap_head **it_frontier = iterated_post_dominance_frontiers(fun,
                                                                set_frontiers, frontiers);
    bool warnings = print_warnings(fun, it_frontier, sets);
    if (!warnings) printf("No potential deadlock found.\n");
    clean_aux_field(fun, 0);
    free_dominance_info(CDI_POST_DOMINATORS);
    return 0;
}
```

Projet d'analyse statique des
fonctions pour détecter les
deadlocks dans les collectives
MPI

Preuve de la slide 7 partie 1

Soit E un ensemble de nœuds

Soit $\text{PDF}(E) = \{ X \mid \exists S \in \text{Succs}(X) \text{ tq } E \gg_{sp} S, E \not\gg_p X \}$

Soit $F = \{ X \mid \exists n \in E \text{ tq } E \gg_{sp} n, X \in \text{PDF}(n), E \not\gg_p X \}$

On va montrer que $\text{PDF}(E) = F$

Soit $x_1 \in F$

$\Leftrightarrow \exists n \in E, x_1 \in \text{PDF}(n), E \gg_p n, E \not\gg_p x_1$

$\Leftrightarrow \exists n \in E, \exists S \in \text{Succ}(x_1), n \gg_{sp} S, n \not\gg_p x_1, E \gg_p n, E \not\gg_p x_1$

$\Rightarrow \exists S \in \text{Succ}(x_1), E \gg_p n \gg_p S, E \not\gg_p x_1$

$\Rightarrow x_1 \in \text{PDF}(E) \Rightarrow F \subseteq \text{PDF}(E)$

Preuve de la slide 7 partie 2

Soit E un ensemble de nœuds

Soit $\text{PDF}(E) = \{ X \mid \exists S \in \text{Succs}(X) \text{ tq } E \gg_{sp} S, E \not\gg_p X \}$

Soit $F = \{ X \mid \exists n \in E \text{ tq } E \gg_{sp} n, X \in \text{PDF}(n), E \not\gg_p X \}$

On va montrer que $\text{PDF}(E) = F$

Soit $x_2 \in \text{PDF}(E)$

$\Leftrightarrow \exists S \in \text{Succ}(x_2), E \gg_{sp} S, E \not\gg_p x_2$

$\Rightarrow S \not\gg_p x_2$, car $E \gg_{sp} S$ et $E \not\gg_p x_2$

donc $x \in \text{PDF}(S)$, car $S \in \text{Succ}(x_2), S \gg_{sp} S$, et $S \not\gg_p x_2$

donc $x_2 \in F \Rightarrow \text{PDF}(E) \subseteq F$

Code lié au warning

```
#pragma Projet_CA mpicoll_check (main, mpi_invalid, mpi_valid)
void mpi_not_checked() {
    MPI_Barrier(MPI_COMM_WORLD);
}
void mpi_invalid(int c) {
    if (c > 5) {
        MPI_Barrier(MPI_COMM_WORLD);
        return;
    } else MPI_Barrier(MPI_COMM_WORLD);

    MPI_Barrier(MPI_COMM_WORLD);
    return;
}
```

```
void mpi_valid(int c) {
    if (c > 5) {
        MPI_Barrier(MPI_COMM_WORLD);
    } else MPI_Barrier(MPI_COMM_WORLD);

    MPI_Barrier(MPI_COMM_WORLD);
    return;
}
int main(int argc, char * argv[])
{
    MPI_Init(&argc, &argv);
    int c = 5;
    if (c < 20) {
        if (c < 10)
        {
            MPI_Barrier(MPI_COMM_WORLD);
        }
        MPI_Barrier(MPI_COMM_WORLD);
    }
    MPI_Finalize();
    return 1;
}
```